

Knowledge Distillation for Data Pruning

Eilis Casey, Zhexi Lu, and Michael Zuo

Group-5

December 5, 2024

Overview

We propose a novel method for data pruning.

Using Knowledge Distillation (KD) methods, we will retain difficult samples for a model to learn.

This will provide a more meaningful dataset for the model to learn on.

Table of Contents

- 1 Introduction
- 2 Related Work
- 3 Methodology
 - Dynamic Data Pruning
 - Training Procedure
- 4 Experiment
 - Experimental Results
- 5 Conclusion

Table of Contents

- 1 Introduction
- 2 Related Work
- 3 Methodology
 - Dynamic Data Pruning
 - Training Procedure
- 4 Experiment
 - Experimental Results
- 5 Conclusion

Introduction

For Large Language Models (LLMs) to achieve such remarkable results, they require a large training corpus

- 570GB of data for GPT-3
- 45TB of data for GPT-4

Thus, training requires substantial computational and financial resources.

Data Pruning

Data Pruning is a technique to reduce the dataset size without compromising performance.

Two types:

- 1 Feature Pruning
- 2 Instance Pruning

Our approach will focus on **instance pruning**.

Data Pruning Methods


Several pruning approaches:

- **Clustering**: Group similar data points and selects representatives
- **Outlier Removal**: Discards anomalous or noisy data points
- **Prototype Selection**: Retain a subset that effectively defines the decision boundary
- **Gradient-based**: Leverage model gradients to evaluate data importance

Our approach focuses on **Prototype Selection** using **Knowledge Distillation** techniques.

Knowledge Distillation

Knowledge Distillation (KD) is a technique for transferring knowledge from a large, complex model (the *teacher*) to a smaller, simpler model (the *student*)



KD Methods

There are several methods for KD:

- 1 **Offline Distillation:** Pretrained teacher model guides training of student
- 2 **Online Distillation:** Teacher and student learn together
- 3 **Self-Distillation:** Special case of online, where deep layers of models help train shallow layers of the same model.

For our approach, we use **offline** KD.

Overview

Overview of our method:

- Fine-tune student model using a pre-trained teacher model
- Examine the Kullback-Leibler (KL) divergence between the teacher and student outputs
- Data points where the KL divergence is larger indicate that the student model has difficulty learning that sample
- Retain these data points, prune the others

Table of Contents

1 Introduction

2 Related Work

3 Methodology

- Dynamic Data Pruning

- Training Procedure

4 Experiment

- Experimental Results

5 Conclusion

Related Work

Another paper “Distilling the knowledge in data pruning” explores using KD with data pruning.

This paper proposes a method for training a student model on a randomly pruned dataset, using KL divergence from a teacher model trained on the full dataset as the key loss component.

Differs from our proposed method because KD has not been used as a tool for pruning but instead as a technique to offset the effects of pruning.

Table of Contents

- 1 Introduction
- 2 Related Work
- 3 Methodology**
 - Dynamic Data Pruning
 - Training Procedure
- 4 Experiment
 - Experimental Results
- 5 Conclusion

Notation

We denote M_t and M_s to be the teacher and student models, respectively.

The models produce logits \mathbf{z}_t and \mathbf{z}_s .

Apply $\text{softmax}(\cdot)$ to the logits to obtain the probability distributions:

$$p_t(y_j|x_i) = \frac{\exp(z_{t,ij})}{\sum_{k=1}^V \exp(z_{t,ik})}$$

$$p_s(y_j|x_i) = \frac{\exp(z_{s,ij})}{\sum_{k=1}^V \exp(z_{s,ik})}$$

Dynamic Data Pruning

To identify informative samples for fine-tuning, compute KL divergence between teacher and student model logits, \mathbf{z}_t and \mathbf{z}_s , respectively.

$$\text{KL}(\mathbf{z}_t, \mathbf{z}_s) = \frac{1}{N} \sum_{i=1}^N \sum_{j=1}^V p_t(y_j|x_i) \log \frac{p_t(y_j|x_i)}{p_s(y_j|x_i)} \quad (1)$$

Where N is the sequence length and V is the vocabulary size

Selection

We select samples where the KL divergence exceeds a threshold τ :

$$\text{Select } x \text{ if } \text{KL}(\mathbf{z}_t, \mathbf{z}_s) > \tau.$$

We increase τ gradually at each epoch, squeezing out fewer, but more divergent, samples in later training steps.

Loss Function

We fine-tune the student model using the selected data samples.

The loss function used is the standard cross-entropy loss for causal language modeling:

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N \log p_s(y_i | y_{<i}; \theta), \quad (2)$$

where θ represents the parameters of the student model, y_i is the target tokens at position i , and $y_{<i}$ is the previously generated tokens before token i .

Fine-tuning with LoRA

To update the student model without incurring high computational costs, we employ Low-Rank Adaptation (LoRA).

LoRA updates weight matrix $\mathbf{W} \in \mathbb{R}^{d \times k}$ by

$$\mathbf{W}' = \mathbf{W} + \Delta\mathbf{W}, \quad \Delta\mathbf{W} = \mathbf{A}\mathbf{B} \quad (3)$$

where $\mathbf{A} \in \mathbb{R}^{d \times r}$ and $\mathbf{B} \in \mathbb{R}^{r \times k}$ are the low-rank matrices with rank $r \ll \min(d, k)$.

Training Algorithm

Algorithm 1 Data Pruning and Fine-Tuning with Dynamic Threshold

Require: Teacher model M_t , student model M_s , dataset \mathcal{D} , initial threshold τ , increase rate α , maximum threshold τ_{\max} , minimum sample count N_{\min}

```
1: Initialize LoRA parameters in  $M_s$ 
2: for each epoch do
3:    $S \leftarrow \{\}$ 
4:   for each sample  $\mathbf{x}$  in  $\mathcal{D}_{\text{train}}$  do
5:     Compute teacher logits  $\mathbf{z}_t$  and student logits  $\mathbf{z}_s$ 
6:     Compute KL divergence  $\text{KL}(\mathbf{z}_t, \mathbf{z}_s)$ 
7:     if  $\text{KL}(\mathbf{z}_t, \mathbf{z}_s) > \tau$  then
8:        $S \leftarrow S \cup \{\mathbf{x}\}$ 
9:     end if
10:  end for
11:  if  $|S| < N_{\min}$  then
12:    Break (Stop training due to insufficient samples)
13:  end if
14:  for each batch in  $S$  do
15:    Compute loss  $\mathcal{L}$  using  $M_s$ 
16:    Backpropagate and update LoRA parameters
17:  end for
18:  Evaluate  $M_s$  on  $\mathcal{D}_{\text{valid}}$ 
19:  Update threshold:  $\tau \leftarrow \min(\tau \cdot (1 + \alpha), \tau_{\max})$ 
20: end for
```

Table of Contents

- 1 Introduction
- 2 Related Work
- 3 Methodology
 - Dynamic Data Pruning
 - Training Procedure
- 4 Experiment
 - Experimental Results
- 5 Conclusion

Experiment Setting

- **Dataset:** We used the first 2000 instruction data in Alpaca. We randomly selected 90% of the data as the training set and 10% of the data as the validation set.
- **Dynamic Data Pruning:** The LoRA configuration is set with $r=8$, lora alpha=16, and a dropout of 0.1.

The AdamW optimizer with a learning rate of 10^{-4} is used to update only the LoRA parameters.

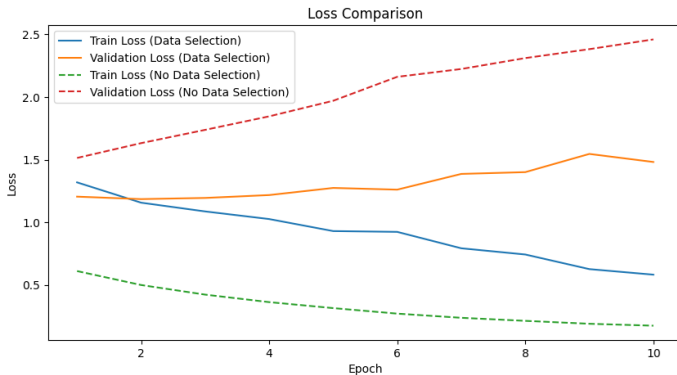
Evaluation Metrics

To assess the performance of the student model, we use the following metrics:

- **Loss:** The average cross-entropy loss over the validation set.
- **Perplexity:** Computed as $\exp(\mathcal{L})$, indicating the model's confidence.
- **Accuracy:** The percentage of correctly predicted tokens or outputs in the validation set.


Experimental Results

Loss



Experimental Results

Perplexity



Experimental Results

Accuracy




Table of Contents

- 1 Introduction
- 2 Related Work
- 3 Methodology
 - Dynamic Data Pruning
 - Training Procedure
- 4 Experiment
 - Experimental Results
- 5 Conclusion

Conclusion

- Dynamic pruning of training samples based on KL divergence improves quality and generalization ability of trained model
- Slower separation of training vs test loss when (over)training without notable loss of accuracy
- Perplexity curve stays relatively flat when applying our data selection technique
- Data is used “less often” throughout training



Thank You!!!